

Exercise 1: Kullback-Leibler Divergence

- (a) You want to approximate the binomial distribution with n number of trials and probability p with a Gaussian distribution with mean μ and variance σ^2 . To find a suitable distribution you investigate the Kullback-Leibler divergence (KLD) in terms of the parameters $\theta = (\mu, \sigma^2)^\top$.
- (i) Write down the KLD for the given setup.
 - (ii) Derive the gradients with respect to θ .
 - (iii) Is there an analytic solution for the optimal parameter setting? If yes, derive the corresponding solution. If no, give a short reasoning.
 - (iv) Independent of the previous exercise, state a numerical procedure to minimize the KLD.
- (b) Sample points according to the true distribution and visualize the KLD for different parameter settings of the Gaussian distribution (including the optimal one if available).
- (c) Create a surface plot with axes n and p and colour value equal to the KLD for the optimal normal distribution.
- (d) Based on the previous result,
- (i) how can the behaviour for varying p be explained?
 - (ii) how can the behaviour for varying n be explained?

Exercise 2: The Convexity of KL Divergence

Let p and q be the PDFs of a pair of absolutely continuous distributions.

- (a) Prove that the KL divergence is convex in the pair (p, q) , i.e.,

$$D_{KL}(\lambda p_1 + (1 - \lambda)p_2 || \lambda q_1 + (1 - \lambda)q_2) \leq \lambda D_{KL}(p_1 || q_1) + (1 - \lambda)D_{KL}(p_2 || q_2), \quad (1)$$

where (p_1, q_1) and (p_2, q_2) are two pairs of distributions and $0 \leq \lambda \leq 1$.

Hint: you can use the log sum inequality, namely that $(a_1 + a_2) \log \left(\frac{a_1 + a_2}{b_1 + b_2} \right) \leq a_1 \log \frac{a_1}{b_1} + a_2 \log \frac{a_2}{b_2}$ holds for $a_1, a_2, b_1, b_2 \geq 0$.