

Solution 1: Kullback-Leibler Divergence

- (a) Let f be the pmf of the $\text{Bin}(n, p)$ distribution and q the density of the $\mathcal{N}(\mu, \sigma^2)$. Recall for $X \sim \text{Bin}(n, p)$, we have $\mathbb{E}_f[X] = np$ and $\text{Var}_f[X] = np(1-p)$. We will first derive a simplified expression for the KL divergence and collect all terms that do not involve μ or σ^2 in a constant.

$$\begin{aligned}
 D_{KL}(f||q) &= \mathbb{E}_f \left[\log \frac{f(X)}{q(X, \theta)} \right] = \underbrace{\mathbb{E}_f[\log f(X)]}_{c_1} - \mathbb{E}_f[\log q(X|\theta)] \\
 &= c_1 - \mathbb{E}_f \left[\log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{1}{2\sigma^2}(X - \mu)^2 \right) \right) \right] \\
 &= c_1 - \mathbb{E}_f \left[\underbrace{-\frac{1}{2} \log(2\pi)}_{c_2} - \frac{1}{2} \log(\sigma^2) - \frac{1}{2\sigma^2}(X - \mu)^2 \right] \\
 &= \underbrace{c_1 + c_2}_{c_3} + \frac{1}{2} \log(\sigma^2) + \frac{1}{2\sigma^2} \mathbb{E}_f[(X - \mu)^2] \\
 &= c_3 + \frac{1}{2} \log(\sigma^2) + \frac{1}{2\sigma^2} \left(\underbrace{\mathbb{E}_f[X^2]}_{\text{Var}_f[X] + \mathbb{E}_f[X]^2} - 2\mu \mathbb{E}_f[X] + \mu^2 \right) \\
 &= c_3 + \frac{1}{2} \log(\sigma^2) + \frac{1}{2\sigma^2} (np(1-p) + (np)^2 - 2\mu np + \mu^2) \\
 &= c_3 + \frac{1}{2} \log(\sigma^2) + \frac{1}{2\sigma^2} (np(1-p) + (np - \mu)^2)
 \end{aligned}$$

- (i) To obtain the gradient of the KLD with respect to the parameter vector $\theta = (\mu, \sigma^2)$, we compute the partial derivatives with respect to the scalar components μ and σ^2 .

$$\frac{\partial D_{KL}(f||q)}{\partial \mu} = \frac{\partial}{\partial \mu} \frac{(np - \mu)^2}{2\sigma^2} = \frac{2(np - \mu) \cdot (-1)}{2\sigma^2} = \frac{\mu - np}{\sigma^2} \quad (1)$$

$$\begin{aligned}
 \frac{\partial D_{KL}(f||q)}{\partial \sigma^2} &= \frac{\partial}{\partial \sigma^2} \left(\frac{1}{2} \log(\sigma^2) + \frac{1}{2\sigma^2} (np(1-p) + (np - \mu)^2) \right) \\
 &= \frac{1}{2\sigma^2} - \frac{1}{2(\sigma^2)^2} (np(1-p) + (np - \mu)^2) \\
 &= \frac{1}{2\sigma^2} - \frac{1}{2\sigma^4} \mathbb{E}_f[(X - \mu)^2]
 \end{aligned} \quad (2)$$

- (ii) Yes, there is. By setting the partial derivatives from before to zero we can identify the critical points. First, we set (1) to zero and get:

$$\frac{\partial D_{KL}(f||q)}{\partial \mu} \stackrel{!}{=} 0 \iff \frac{\hat{\mu} - np}{\hat{\sigma}^2} = 0 \iff \hat{\mu} = np = \mathbb{E}_f[X]$$

Next, we set equation (2) to zero and plug in our solution for $\hat{\mu}$:

$$\begin{aligned}
 \frac{\partial D_{KL}(f||q)}{\partial \sigma^2} \stackrel{!}{=} 0 &\iff \frac{1}{2\hat{\sigma}^2} = \frac{1}{2\hat{\sigma}^4} \mathbb{E}_f[(X - \hat{\mu})^2] \quad | \cdot 2\hat{\sigma}^4 \\
 &\iff \hat{\sigma}^2 = \mathbb{E}_f[(X - \mathbb{E}_f[X])^2] = \text{Var}_f[X] = np(1-p)
 \end{aligned}$$


Note that in order to prove that the critical point is a minimum, we would have to check that the Hessian of $D_{KL}(f||q)$ with respect to (μ, σ^2) is positive definite at $(\hat{\mu}, \hat{\sigma}^2)$.

(iii) We could, alternatively, use the gradients and do gradient descent to find the optimal θ .

```
(b) nr_points = 1000
p = 0.5
n = 100
# create data
X <- rbinom(nr_points, prob = p, size = n)

# define different Normal density functions
normal_optimal <- function(x) dnorm(x, mean = n*p, sd = sqrt(n*p*(1-p)))
normal_shift <- function(x) dnorm(x, mean = n*p - 10, sd = sqrt(n*p*(1-p)))
normal_scale_increase <- function(x) dnorm(x, mean = n*p, sd = sqrt(n*p*(1-p))*2)
normal_right_scale_decrease <- function(x) dnorm(x, mean = n*p + 20, sd = p*(1-p))

hist(X, breaks = 25, xlim = c(10, 100), freq = FALSE)
curve(normal_optimal, from = 10, to = 100, add = TRUE, col = "green")
curve(normal_shift, from = 10, to = 100, add = TRUE, col = "blue")
curve(normal_scale_increase, from = 10, to = 100, add = TRUE, col = "orange")
curve(normal_right_scale_decrease, from = 10, to = 100, add = TRUE, col = "red")
```



For these distributions, we get the following KL divergence values (up to an additive constant):

$$D_{KL}(f||q) = c_3 + 0.5 \log \sigma^2 + \frac{1}{2\sigma^2} (\text{Var}_f(X) + (np - \mu)^2)$$

```
kld_value <- function(mu,sigma2)
{
  0.5*log(sigma2) +
  0.5 * (sigma2)^(-1) * (n*p*(1-p) + (n*p - mu)^2)
}
(optimal_green <- kld_value(n*p,n*p*(1-p)))

## [1] 2.109438

(shift_blue <- kld_value(n*p-10,n*p*(1-p)))
```

```
## [1] 4.109438

(scale_increase_orange <- kld_value(n*p,n*p*(1-p)*4))

## [1] 2.427585

(right_scale_decrease_red <- kld_value(n*p+20, (p*(1-p))^2))

## [1] 3398.614
```

- (c) Since we are now required to calculate the exact KLD values, we would also have to calculate $\mathbb{E}_f(\log f(X))$, which is somewhat more difficult. If you search the internet for a solution (\rightarrow “entropy of a binomial distribution”), you will find an approximate solution using the de-Moivre-Laplace theorem. Alternatively, we could make use of the central limit theorem, but then we would just approximate f with a normal distribution with $\mu = np$ and $\sigma^2 = np(1 - p)$, which would give us a constant KLD of zero (the very same happens if you use the first approximation using the de-Moivre-Laplace-theorem). We here instead will approximate the expectation using a large sample from the true underlying distribution:

$$D_{KL}(f||q) \approx \frac{1}{B} \sum_{b=1}^B [\log f(X) - \log q(X|\mu = np, \sigma^2 = np(1 - p))]$$

```
p_seq <- seq(0.01, 0.99, l = 100)
n_seq <- seq(10, 500, by = 100)
B <- 10000

kld_value_approx <- function(n,p){


  # sample a large number of data points from true distribution
  x <- rbinom(B, prob = p, size = n)

  # approximate the mean; threshold values to 0 if < 0 due
  # to the approximation
  pmax(
    mean(
      dbinom(x, prob = p, size = n, log = TRUE) -
      dnorm(x, mean = n*p, sd = sqrt(n*p*(1-p)), log = TRUE),
      na.rm = TRUE
    ),
    0)
}

kld_val <- sapply(n_seq, function(this_n)
  sapply(p_seq, function(this_p) kld_value_approx(this_n, this_p)))

cols = rev(colorRampPalette(c('darkred','red','blue','lightblue'))(50))

filled.contour(x = p_seq, y = n_seq, z = kld_val,
  xlab = "p", ylab = "n",
  col = cols
)
```



- (d) Based on the previous result, one can see that the KLD is very close to zero but has larger values for very small or very large values of p and n in combination with a small number of experiments n . These are exactly the cases where the normal approximation of a binomial distribution does not work so well.

Solution 2: The Convexity of KL Divergence

- (a) We expand the left side of the inequality and obtain:

$$\begin{aligned}
 & D_{KL}(\lambda p_1 + (1 - \lambda)p_2 \| \lambda q_1 + (1 - \lambda)q_2) \\
 &= \int_{\mathcal{X}} \left((\lambda p_1(x) + (1 - \lambda)p_2(x)) \log \frac{\lambda p_1(x) + (1 - \lambda)p_2(x)}{\lambda q_1(x) + (1 - \lambda)q_2(x)} \right) dx \\
 &\leq \int_{\mathcal{X}} \left(\lambda p_1(x) \log \frac{p_1(x)}{q_1(x)} + (1 - \lambda)p_2(x) \log \frac{(1 - \lambda)p_2(x)}{(1 - \lambda)q_2(x)} \right) dx \\
 &= \lambda \int_{\mathcal{X}} \left(p_1(x) \log \frac{p_1(x)}{q_1(x)} \right) dx + (1 - \lambda) \int_{\mathcal{X}} \left(p_2(x) \log \frac{p_2(x)}{q_2(x)} \right) dx \\
 &= \lambda D_{KL}(p_1 \| q_1) + (1 - \lambda) D_{KL}(p_2 \| q_2).
 \end{aligned} \tag{3}$$