


Supervised Learning

Practical Tips for Feature Selection



Learning goals

- Add learning goals

SOLVING A FEATURE SELECTION PROBLEM (TAKEN FROM GUYON (2003))

❶ Do you have domain knowledge?

If yes, construct a better set of “ad hoc” features.

❷ Are your features commensurate?

If no, consider normalizing them.

❸ Do you suspect interdependence of features?

If yes, expand your feature set by constructing conjunctive features or products of features, as many as your computer resources allow you to.

❹ Do you need to prune the input features (e.g. for cost, speed or data-understanding reasons)?

If no, construct disjunctive features or weighted sums of features (e.g. by clustering or matrix factorization).

SOLVING A FEATURE SELECTION PROBLEM (TAKEN FROM GUYON (2003))

- 5 Do you need to assess features individually (e.g. to understand their influence on the system, or because their number is so large that you need to do a first filtering)?**
If yes, use a variable-ranking method. Otherwise, do it anyway to get baseline results.
- 6 Do you need a predictor?**
If no, stop.
- 7 Do you suspect your data is “dirty” (has a few meaningless input patterns and/or noisy outputs or wrong class labels)?**
If yes, detect the outlier examples using the top-ranking features obtained in step 5 as representation; check and/or discard them.

SOLVING A FEATURE SELECTION PROBLEM (TAKEN FROM GUYON (2003))

8 Do you know what to try first?

If no, use a linear predictor. Use a forward selection method with the “probe” method as a stopping criterion or use the L_0 norm embedded method. For comparison, following the ranking of step 5, construct a sequence of predictors from the same family, using increasing subsets of features. Can you match or improve performance with a smaller subset? If yes, try a nonlinear predictor with that subset.

9 Do you have new ideas, time, computational resources, and enough examples?

If yes, compare several feature selection methods, including your new idea, correlation coefficients, backward selection and embedded methods. Use linear and nonlinear predictors. Select the best approach via model selection.

SOLVING A FEATURE SELECTION PROBLEM (TAKEN FROM GUYON (2003))

- 10 Do you want a stable solution (to improve performance and/or understanding)?

If yes, sub-sample your data and redo your analysis for several bootstraps.

OPEN POINTS / PROBLEMS

- In general, it is difficult to give suggestions on when to use which feature selection method.
- Most of the time, it is reasonable to start with a simple, fast method. If this yields unsatisfactory results, one can gradually move to more expensive methods.
- Not every introduced method can be generalized to multi-class problems in an easy fashion.
- Combining the choice of an appropriate classifier and parameter tuning with feature selection is not simple.