# Introduction to Machine Learning

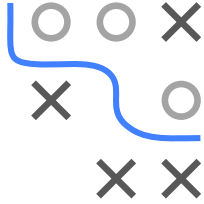

## Regularization
## Bias-variance Tradeoff

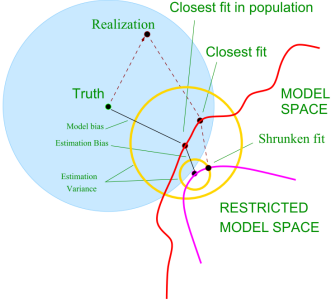


**Learning goals**

- Understand the bias-variance trade-off
- Know the definition of model bias, estimation bias, and estimation variance

# BIAS-VARIANCE TRADEOFF

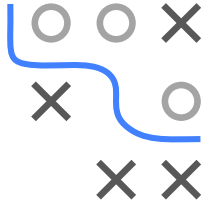

In this slide set, we will visualize the bias-variance trade-off.

First, we start with a DGP $\mathbb{P}_{xy}$ and a suitable loss function $L : \mathbb{R}^g \times \mathbb{R}^g \to \mathbb{R}$ where $\mathbb{R}^g$ is numerical encoding of $\mathcal{Y}$. We measure the distance between models $f : \mathcal{X} \to \mathbb{R}^g$ via

$$d(f, f') = \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_{\mathbf{x}}} \left[ L(f(\mathbf{x}), f'(\mathbf{x}) \right] .$$

We restrict our attention to losses for which $d$ becomes a metric, e.g., L1-loss, L2-loss, etc.

We define $f_{\text{true}}$ as the risk minimizer such that

$$f_{\text{true}} \in \arg\min_{f \in \mathcal{H}_0} \mathbb{E}_{(\mathbf{x},y) \sim \mathbb{P}_{xy}} \left[ L(y, f(\mathbf{x})) \right]$$

where $\mathcal{H}_0 = \{ f : \mathcal{X} \to \mathbb{R}^g | \; d(\underline{0}, f) < \infty \}$ and $\underline{0} : \mathcal{X} \to \{0\}$.

# BIAS-VARIANCE TRADEOFF / 2

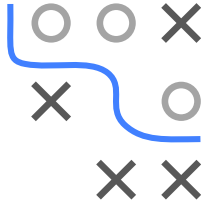

In practice, our model space $\mathcal{H}$ usually is a proper subset of $\mathcal{H}_0$ and in general $f_{\text{true}} \notin \mathcal{H}$.
We define $f^*$ as the risk minimizer in $\mathcal{H}$, i.e.,

$$f^* \in \arg\min_{f \in \mathcal{H}} \mathbb{E}_{(\mathbf{x},y)\sim\mathbb{P}_{xy}} \left[ L(f(\mathbf{x}, y)) \right].$$

It is the function in $\mathcal{H}$ closest to $f_{\text{true}}$, and we call $d(f_{\text{true}}, f^*)$ the model bias.

# BIAS-VARIANCE TRADEOFF / 3

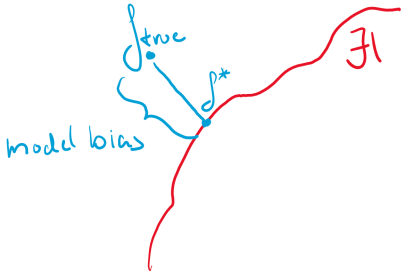


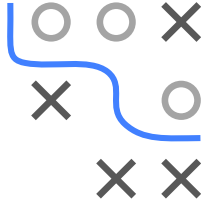

# BIAS-VARIANCE TRADEOFF / 4

We can further restrict the model space such that $\mathcal{H}_R$ is a proper subset of $\mathcal{H}$. We define $f_R^*$ as the risk minimizer in $\mathcal{H}_R$, i.e.,

$$f_R^* \in \arg\min_{f \in \mathcal{H}_R} \mathbb{E}_{(\mathbf{x},y)\sim\mathbb{P}_{xy}} \left[ L(f(\mathbf{x}), y) \right].$$

It is the function in $\mathcal{H}_R$ closest to $f_{\text{true}}$, and we call $d(f_R^*, f^*)$ the estimation bias.

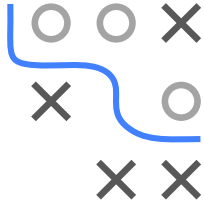

# BIAS-VARIANCE TRADEOFF / 5

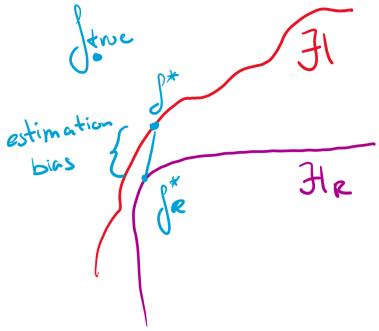


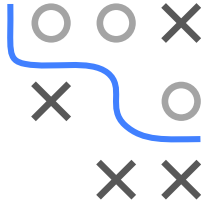

# BIAS-VARIANCE TRADEOFF / 6

We sample a finite dataset $\mathcal{D} = \left(\mathbf{x}^{(i)}, y^{(i)}\right)^n \in (\mathbb{P}_{xy})^n$ and find via ERM

$$\hat{f} \in \arg\min_{f \in \mathcal{H}} \sum_{i=1}^{n} L\left(y^{(i)}, \hat{f}(\mathbf{x}^{(i)})\right).$$

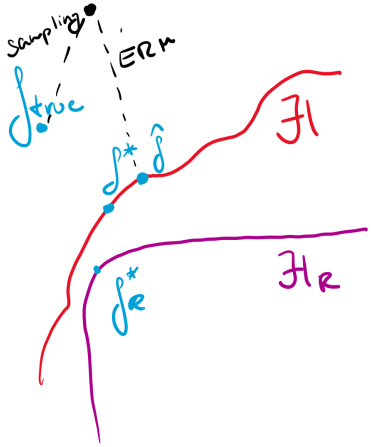


Note:

- $L : \mathcal{Y} \times \mathbb{R}^g \to \mathbb{R}$ is overloaded.
- The samples are only shown in the visualization for didactic purposes but are not an element of $\mathcal{H}$.

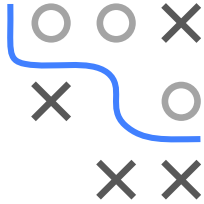

# BIAS-VARIANCE TRADEOFF / 7

Let's assume that $\hat{f}$ is an unbiased estimate of $f^*$ (e.g., valid for linear regression), and we repeat the sampling process of $\hat{f}$.

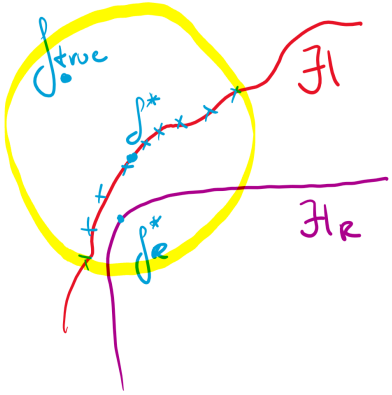


- We can measure the spread of sampled $\hat{f}$ around $f^*$ via $\delta = \mathrm{Var}_{\mathcal{D}}\left[d(f^*, \hat{f})\right]$ which we call the estimation variance.
- We visualize this as a circle around $f^*$ with radius $\delta$.

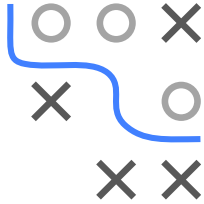

# BIAS-VARIANCE TRADEOFF / 8

We repeat the previous construction in the restricted model space $\mathcal{H}_R$ and sample $\hat{f}_R$ such that

$$\hat{f}_R \in \arg\min_{f \in \mathcal{H}_R} \sum_{i=1}^{n} L\left(y^{(i)}, \hat{f}(\mathbf{x}^{(i)})\right).$$

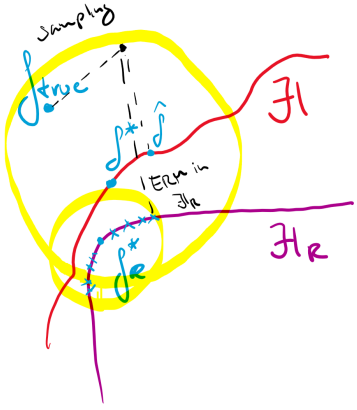


- We can measure the spread of sampled $\hat{f}_R$ around $f_R^*$ via $\delta = \mathrm{Var}_{\mathcal{D}}\left[d(f^*, \hat{f}_R)\right]$ which we also call estimation variance.
- We observe that the increased bias results in a smaller estimation variance in $\mathcal{H}_R$ compared to $\mathcal{H}$.